

# PCA v/s LDA v/s SVD

Riya Goyal 2148054

## Importing libraries

```
In [2]: 1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import math
6 from sklearn.model_selection import train_test_split
```

## Data Description:

- The dataset contains transactions made by credit cards in September 2013 by European cardholders.
- This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.
- Feature **'Time'** contains the seconds elapsed between each transaction and the first transaction in the dataset.
- Feature **'Amount'** is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning.
- Feature **'Class'** is the response variable and it takes value 1 in case of fraud and 0 otherwise.

```
In [3]: 1 # importing the dataset
2 df = pd.read_csv('D:/Downloads/creditcard.csv/creditcard.csv')
3
4 # displaying the first 5 rows of the dataset
5 df.head()
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458

5 rows × 31 columns

```
In [4]: 1 df.shape
```

```
Out[4]: (284807, 31)
```

## Dropping unnecessary columns, which contains null values

```
In [5]: 1 # columns containing null values
2 df.columns[df.isnull().any()]
```

```
Out[5]: Index([], dtype='object')
```

There is no missing data in the dataset

```
In [6]: 1 df['Class'].value_counts()
```

```
Out[6]: 0    284315
1         492
Name: Class, dtype: int64
```

There is strong imbalance in the target value counts. To solve this issue, we will oversample the dataset.

## Oversampling the dataset

```
In [7]: 1 def _oversample_positives(df, target):
2       """ Oversample the minority classes to match
3       the majority class.
4
5       :param df: pandas dataframe - input df.
6       :param target: string - classification target column.
7
8       :return: pandas dataframe - oversampled version
9       """
10
11       class_count = df[target].value_counts()
12
13       print("Before oversampling: %s" % class_count)
14
15       for i in range(1,len(class_count)):
16           df_i = df[df[target] == i]
17           oversampling_factor_i = class_count[0] / float(class_count[i])
18           print(len(df_i))
19           print("Oversampling factor for class %i: %s" %(i, str(oversampling_factor_i)))
20
21           # Integer part of oversampling
22           df = df.append(
23               [df_i] * int(math.floor(oversampling_factor_i) - 1),
24               ignore_index=False)
25
26           # Float part of oversampling
27           df = df.append(
28               [df_i.sample(frac=oversampling_factor_i % 1)],
29               ignore_index=False)
30
31       print("After oversampling: %s" % df[target].value_counts())
32       print("Shape after oversampling: %s" % str(df.shape))
33
34       return df
```

```
In [8]: 1 df_oversampled = _oversample_positives(df, 'Class')
```

```
Before oversampling: 0    284315
1         492
Name: Class, dtype: int64
492
Oversampling factor for class 1: 577.8760162601626
After oversampling: 0    284315
1    284315
Name: Class, dtype: int64
Shape after oversampling: (568630, 31)
```

## Splitting the dataset into Independent and Dependent variables

```
In [9]: 1 X = df_oversampled.drop(['Class'],axis = 1)
2       y = df_oversampled['Class']
```

```
In [10]: 1 y.value_counts()
```

```
Out[10]: 0    284315
1    284315
Name: Class, dtype: int64
```

## Training data and testing data

```
In [11]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=1)
```

## Standardizing the data

```
In [12]: 1 from sklearn.preprocessing import StandardScaler
2       sc = StandardScaler()
3       X_train = sc.fit_transform(X_train)
4       X_test = sc.transform(X_test)
```

## Linear Discriminant Analysis

Linear Discriminant Analysis as its name suggests is a linear model for classification and dimensionality reduction. Most commonly used for feature extraction in pattern classification problems.

### Assumptions:

LDA makes some assumptions about the data:

- Assumes the data to be distributed normally or Gaussian distribution of data points i.e. each feature must make a bell-shaped curve when plotted.
- Each of the classes has identical covariance matrices.

However, it is worth mentioning that LDA performs quite well even if the assumptions are violated.

```
In [13]: ▶ 1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
2 lda = LDA(n_components=1) #n_components to be less than the n_classes count
3 X_train = lda.fit_transform(X_train, y_train)
4 X_test = lda.transform(X_test)
5 LDA_df = pd.DataFrame(X_test)
6 LDA_df.head(10)
```

```
Out[13]:
      0
0  1.756693
1  1.403826
2 -0.689900
3 -2.580457
4  1.359025
5 -0.806420
6  0.152485
7  0.525899
8 -0.154159
9  0.276191
```

## Modelling - Logistic Regression

```
In [14]: ▶ 1 from sklearn.linear_model import LogisticRegression
2 classifier = LogisticRegression(random_state=0)
3 classifier.fit(X_train, y_train)
4 y_pred = classifier.predict(X_test)
5 y_pred2 = classifier.predict(X_train)
```

## Model Evaluation

```
In [15]: ▶ 1 from sklearn.metrics import confusion_matrix, accuracy_score
2
3 cm = confusion_matrix(y_test, y_pred)
4
5 print(cm)
6 print('Accuracy: ' + str(accuracy_score(y_test, y_pred)))
```

```
[[54629  2297]
 [ 5740 51060]]
Accuracy: 0.9293301443821114
```

## Principal Component Analysis

Principal component analysis (PCA) is a technique for reducing the dimensionality of such datasets, increasing interpretability but at the same time minimizing information loss. It does so by creating new uncorrelated variables that successively maximize variance.

```
In [16]: 1 # Importing standardscalar module
2 from sklearn.preprocessing import StandardScaler
3 df = pd.DataFrame(X)
4 y_target = pd.DataFrame(y)
5 scalar = StandardScaler()
6
7 # fitting
8 scalar.fit(df)
9 scaled_data = scalar.transform(df)
10
11 # Importing PCA
12 from sklearn.decomposition import PCA
13
14 # Let's say, components = 2
15 pca = PCA(n_components = 0.95)
16 pca.fit(scaled_data)
17 x_pca = pca.transform(scaled_data)
18 print(x_pca.shape)
19 PCA_df = pd.DataFrame(x_pca)
20 PCA_df.head()
```


(568630, 17)

Out[16]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	-2.017422	-0.079591	0.270550	0.674383	-0.385741	1.409463	-1.026928	0.507152	-1.805759	-0.222482	0.500370	0.295596	-0.03115
1	-1.938631	-0.518955	0.009032	-0.092563	-0.703678	0.711102	-1.065389	-0.913978	-1.023532	-0.655152	0.747481	-0.203345	0.08095
2	-2.099587	0.313580	0.095831	0.335490	0.099113	2.515647	-0.768611	-0.864740	-1.587526	-1.435280	0.386839	-0.321086	1.40594
3	-2.340797	-0.038814	-0.452258	-1.047547	-1.109917	1.074414	0.101514	-1.381527	-0.728062	-1.592945	0.695028	0.530310	-0.65105
4	-1.898936	-0.053823	0.790057	0.433809	-0.245760	0.886058	-1.582719	-0.982657	0.132393	-0.345936	1.300697	-0.919169	-0.41544

17 variables justify 95% variability in my data

```
In [17]: 1 plt.plot(np.cumsum(pca.explained_variance_ratio_))
2 plt.xlabel('number of components')
3 plt.ylabel('cumulative explained variance');
```



## Training data and testing data

```
In [18]: 1 X1_train, X1_test, y1_train, y1_test = train_test_split(PCA_df, y, test_size=0.2, random_state=1)
```

## Logistic Regression Model

```
In [19]: 1 from sklearn.linear_model import LogisticRegression
2 classifier = LogisticRegression(random_state=0)
3 classifier.fit(X1_train, y1_train)
4 y1_pred = classifier.predict(X1_test)
5 y1_pred2 = classifier.predict(X1_train)
```

## Model Evaluation

```
In [20]: 1 from sklearn.metrics import confusion_matrix, accuracy_score
2
3 cm1 = confusion_matrix(y1_test, y1_pred)
4
5 print(cm1)
6 print('Accuracy: ' + str(accuracy_score(y1_test, y1_pred)))
```

```
[[55550 1376]
 [ 5481 51319]]
Accuracy: 0.9397059599388002
```

## Inference:

LDA & PCA both performed similarly on the dataset contributing to the model accuracy which was observed to be **92.93%** and **93.96%** respectively. The **information contained within 60 features was retained in 17 features approximately** and we did not lose much of information as we saw good accuracy.

```
In [21]: 1 from sklearn.preprocessing import StandardScaler
2 sc = StandardScaler()
3
4 #fitting
5 sc.fit(X)
6 X = sc.transform(X)
```

```
In [22]: 1 from sklearn.decomposition import TruncatedSVD
2
3 print("Original Matrix:")
4 print(X, '\n')
5
6 svd = TruncatedSVD(n_components = 4)
7 X_transf = svd.fit_transform(X)
8
9 print("Singular values: \n")
10 print(svd.singular_values_, '\n')
11
12 print("Transformed Matrix after reducing to 2 features: \n")
13 print(X_transf)
```

Original Matrix:

```
[[-1.82317125  0.18495473 -0.50611358 ...  0.04759203 -0.12984457
  0.17479509]
 [-1.82317125  0.64671985 -0.41494358 ... -0.09265743 -0.05086622
 -0.40402063]
 [-1.82315048  0.18521769 -0.84702927 ... -0.13828151 -0.21527228
  1.07707476]
 ...
 [ 1.09027437  0.60313741  0.38847542 ...  0.08311508  0.1556491
 -0.41158427]
 [-0.35229583  0.46663122 -0.16650073 ...  0.04922196  0.48465339
 -0.41162367]
 [ 1.30211714  0.76634433 -0.19882401 ... -0.04222082  0.00619462
 -0.41067821]]
```

Singular values:

```
[2542.83877106 1295.96156021 1090.67602695  966.24091486]
```

Transformed Matrix after reducing to 2 features:

```
[[-2.01742172 -0.07958723  0.27057454  0.67440123]
 [-1.93863053 -0.518954   0.00904835 -0.09269215]
 [-2.09958693  0.31359854  0.0959183   0.33562026]
 ...
 [-1.82230282 -0.11685556 -1.42931602 -2.3445905 ]
 [-1.96656279  0.29998635 -0.47543002 -0.77941364]
 [-2.24188402  0.52321109 -0.995532   -2.0965485  ]]
```

```
In [23]: 1 X_train, X_test, y_train, y_test = train_test_split(X_transf, y, test_size=0.2, random_state=1)
```

```
In [24]: 1 from sklearn.linear_model import LogisticRegression
2 classifier = LogisticRegression(random_state=1)
3 classifier.fit(X_train, y_train)
4 y_pred = classifier.predict(X_test)
5 y_pred2 = classifier.predict(X_train)
```

```
In [25]: 1 from sklearn.metrics import confusion_matrix, accuracy_score
2
3 cm = confusion_matrix(y_test, y_pred)
4
5 print(cm)
6 print('Accuracy: ' + str(accuracy_score(y_test, y_pred)))
```

```
[[56398  528]
 [ 9427 47373]]
Accuracy: 0.9124650475704764
```

```
In [32]: 1 var_explained = svd.explained_variance_ratio_
2 var_explained
```

```
Out[32]: array([0.37904138, 0.09845397, 0.06973335, 0.05472931])
```

```
In [33]: 1 def select_n_components(var_ratio, goal_var: float) -> int:
2     # Set initial variance explained so far
3     total_variance = 0.0
4
5     # Set initial number of features
6     n_components = 0
7
8     # For the explained variance of each feature:
9     for explained_variance in var_ratio:
10
11         # Add the explained variance to the total
12         total_variance += explained_variance
13
14         # Add one to the number of components
15         n_components += 1
16
17         # If we reach our goal level of explained variance
18         if total_variance >= goal_var:
19             # End the loop
20             break
21
22     # Return the number of components
23     return n_components
24 # Run function
```

```
In [34]: 1 select_n_components(var_explained, 0.95)
```

```
Out[34]: 4
```

The model performance with **SVD** had an accuracy of **91.24%** since we scaled down the whole dataset into 4 features, while PCA gave us an accuracy of **93.96%** with 17 components. **LDA** performed with accuracy of **92.93%**. With all the three techniques, we did not witness any major loss of information of our original data but we can say that LDA and SVD work better than PCA, given the number of features in each technique.

```
In [ ]: 1
```